

Class 09: Candy Mini Project

Kris Price (PID: A17464127)

Table of contents

Importing Candy Data	1
What is in the dataset?	2
What is your favorite candy?	2
Exploratory Analysis	4
Overall Candy Rankings	6
Time to add some useful color	9
Taking a look at pricepercent	10
Exploring the correlation structure	12
Principal Component Analysis	13
Summary	16
Optional extension questions	16

Importing Candy Data

Our dataset is a CSV file so we use `read.csv()`

```
candy <- read.csv("candy-data.csv", row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat	crisp	pedricewafer
100 Grand	1	0	1		0	0		1
3 Musketeers	1	0	0		0	1		0
One dime	0	0	0		0	0		0
One quarter	0	0	0		0	0		0
Air Heads	0	1	0		0	0		0
Almond Joy	1	0	0		1	0		0
	hard	bar	pluribus	sugarpercent	pricepercent	winpercent		
100 Grand	0	1	0	0.732	0.860	66.97173		

3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

What is in the dataset?

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

[1] 85

There are 85 candies.

Q2. How many fruity candy types are in the dataset?

```
table(candy$fruity)
```

0	1
47	38

There are 38 fruity candies.

What is your favorite candy?

Q3. What is your favorite candy (other than Twix) in the dataset and what is its winpercent value?

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
candy %>%
  filter(row.names(candy) == "Sour Patch Kids") %>%
  select(winpercent)
```

```
winpercent
Sour Patch Kids      59.864
```

Sour Patch Kids have a winpercent value of 59.864.

Q4. What is the winpercent value for “Kit Kat”?

```
candy %>%
  filter(row.names(candy) == "Kit Kat") %>%
  select(winpercent)
```

```
winpercent
Kit Kat      76.7686
```

Kit Kats have a winpercent value of 76.7686.

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy %>%
  filter(row.names(candy) == "Tootsie Roll Snack Bars") %>%
  select(winpercent)
```

```
winpercent
Tootsie Roll Snack Bars    49.6535
```

Tootsie Roll Snack Bars have a winpercent value of 49.6535.

There is a useful `skim()` function in the `skimr` package that can help give you a quick overview of a given dataset. Let’s install this package using `install.packages("skimr")` and try it on our candy data.

```
#library(skimr)
#skim(candy)
```

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The `winpercent` variable is the only one that's not on a scale from 0-1.

Q7. What do you think a zero and one represent for the `candy$chocolate` column?

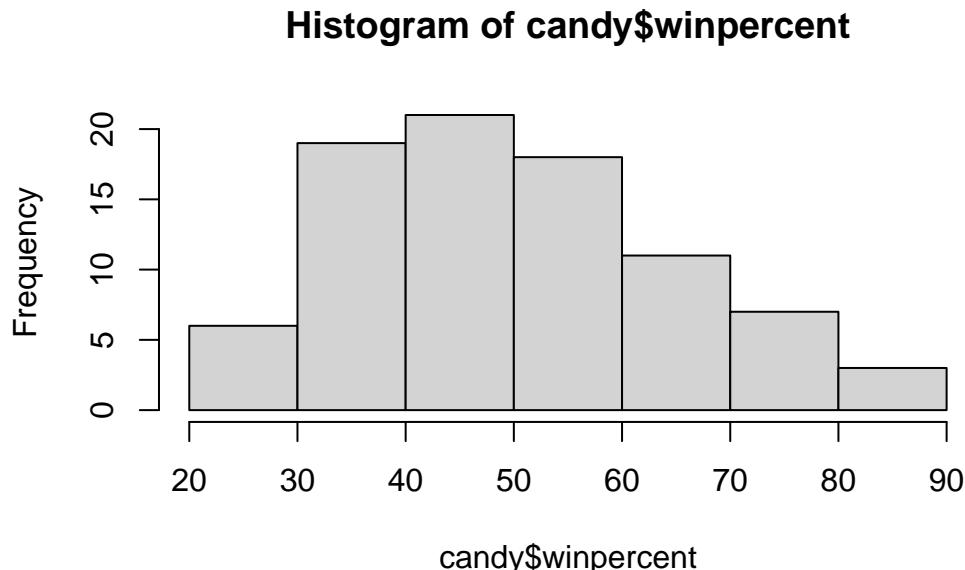
The `candy$chocolate` column has binary values: 0 means that the candy isn't chocolate, while 1 means that the candy is chocolate.

Exploratory Analysis

Q8. Plot a histogram of `winpercent` values

Using base R:

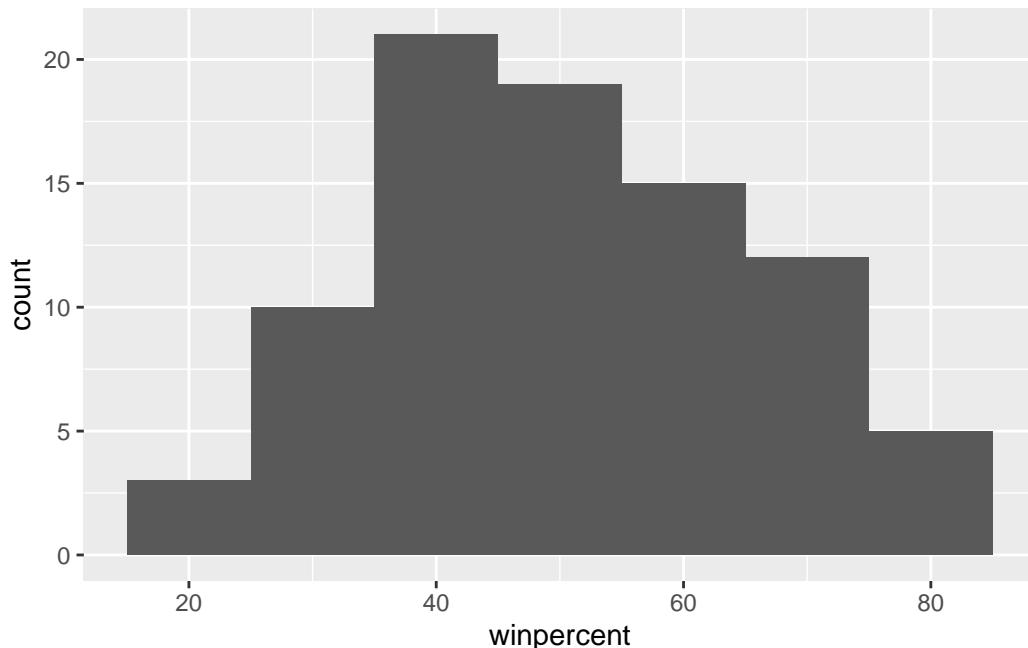
```
hist(candy$winpercent)
```



Using ggplot2:

```
library(ggplot2)

ggplot(candy, aes(x = winpercent)) +
  geom_histogram(binwidth = 10)
```



Q9. Is the distribution of winpercent values symmetrical?

The distribution is not symmetrical (it's right-skewed).

Q10. Is the center of the distribution above or below 50%?

```
summary(candy$winpercent)
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	22.45	39.14	47.83	50.32	59.86	84.18

The center of this distribution is below 50% (median = 47.83).

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
candy_c <- candy %>%
  filter(chocolate == 1)
summary(candy_c$winpercent)
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.  
34.72 50.35 60.80 60.92 70.74 84.18
```

```
candy_f <- candy %>%  
  filter(fruity == 1)  
summary(candy_f$winpercent)
```

```
Min. 1st Qu. Median Mean 3rd Qu. Max.  
22.45 39.04 42.97 44.12 52.11 67.04
```

On average, chocolate candy (mean = 60.92) is higher-ranked than fruit candy (mean 44.12).

Q12. Is this difference statistically significant?

```
t.test(candy_c$winpercent, candy_f$winpercent)
```

Welch Two Sample t-test

```
data: candy_c$winpercent and candy_f$winpercent  
t = 6.2582, df = 68.882, p-value = 2.871e-08  
alternative hypothesis: true difference in means is not equal to 0  
95 percent confidence interval:  
 11.44563 22.15795  
sample estimates:  
mean of x mean of y  
60.92153 44.11974
```

Yes, this difference is statistically significant ($p < 0.05$).

Overall Candy Rankings

Q13. What are the five least liked candy types in this set?

```
low <- candy %>%  
  arrange(winpercent)  
head(low, 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0	0	0	0
Boston Baked Beans	0	0	0	1	0	0
Chiclets	0	1	0	0	0	0
Super Bubble	0	1	0	0	0	0
Jawbusters	0	1	0	0	0	0
	crisped	rice	wafer	hard	bar	pluribus
	sugar	percent	price	percent	win	percent
Nik L Nip	0	0	0	1	0.197	0.976
Boston Baked Beans	0	0	0	1	0.313	0.511
Chiclets	0	0	0	1	0.046	0.325
Super Bubble	0	0	0	0	0.162	0.116
Jawbusters	0	1	0	1	0.093	0.511
	win	percent				
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble, Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?

```
high <- candy %>%
  arrange(desc(winpercent))

head(high, 5)
```

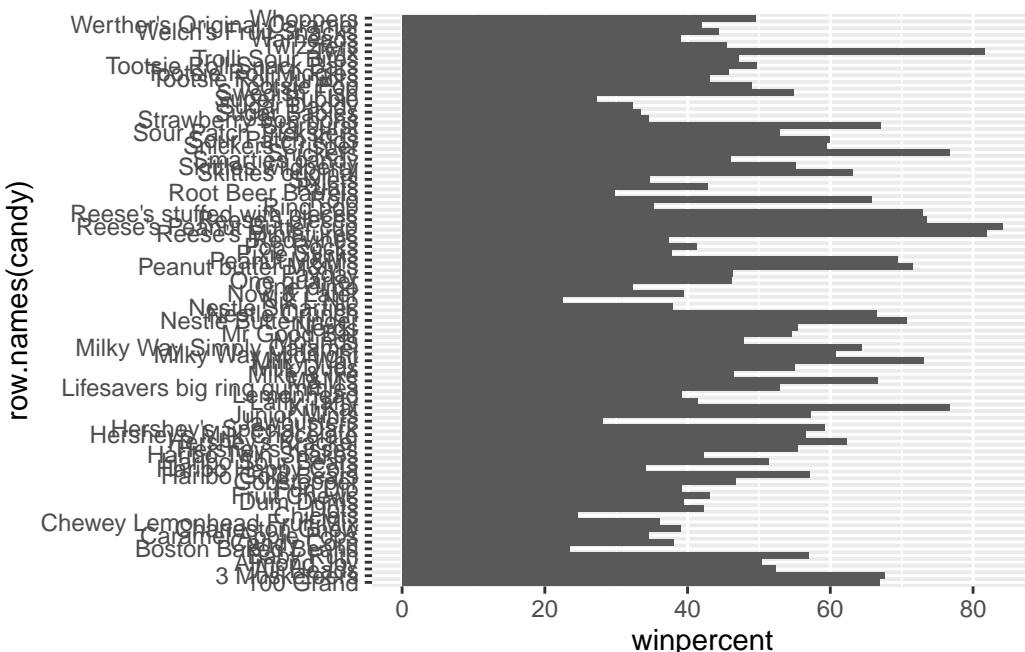
	chocolate	fruity	caramel	peanuty	almondy	nougat
Reese's Peanut Butter cup	1	0	0	1	0	0
Reese's Miniatures	1	0	0	1	0	0
Twix	1	0	1	0	0	0
Kit Kat	1	0	0	0	0	0
Snickers	1	0	1	1	1	1
	crisped	rice	wafer	hard	bar	pluribus
	sugar	percent	price	percent	win	percent
Reese's Peanut Butter cup	0	0	0	0	0.720	
Reese's Miniatures	0	0	0	0	0.034	
Twix	1	0	1	0	0.546	
Kit Kat	1	0	1	0	0.313	
Snickers	0	0	1	0	0.546	
	price	percent	win	percent		
Reese's Peanut Butter cup	0.651	84.18029				

Reese's Miniatures	0.279	81.86626
Twix	0.906	81.64291
Kit Kat	0.511	76.76860
Snickers	0.651	76.67378

Reese's Peanut Butter cup, Reese's Miniatures, Twix, Kit Kat, Snickers.

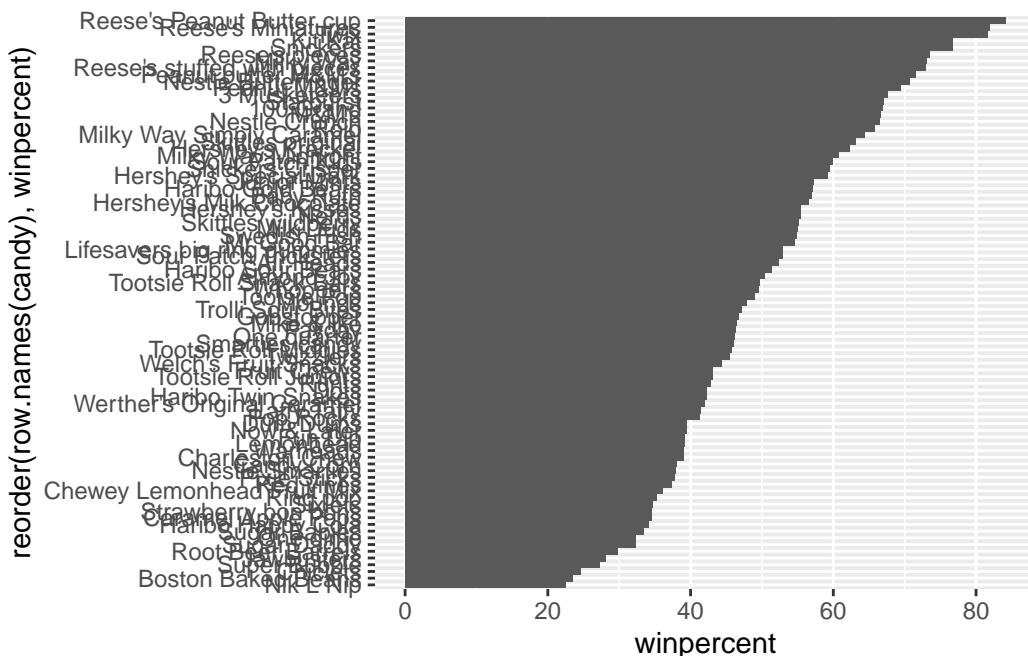
Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy, aes(winpercent, row.names(candy))) +
  geom_col()
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by `winpercent`?

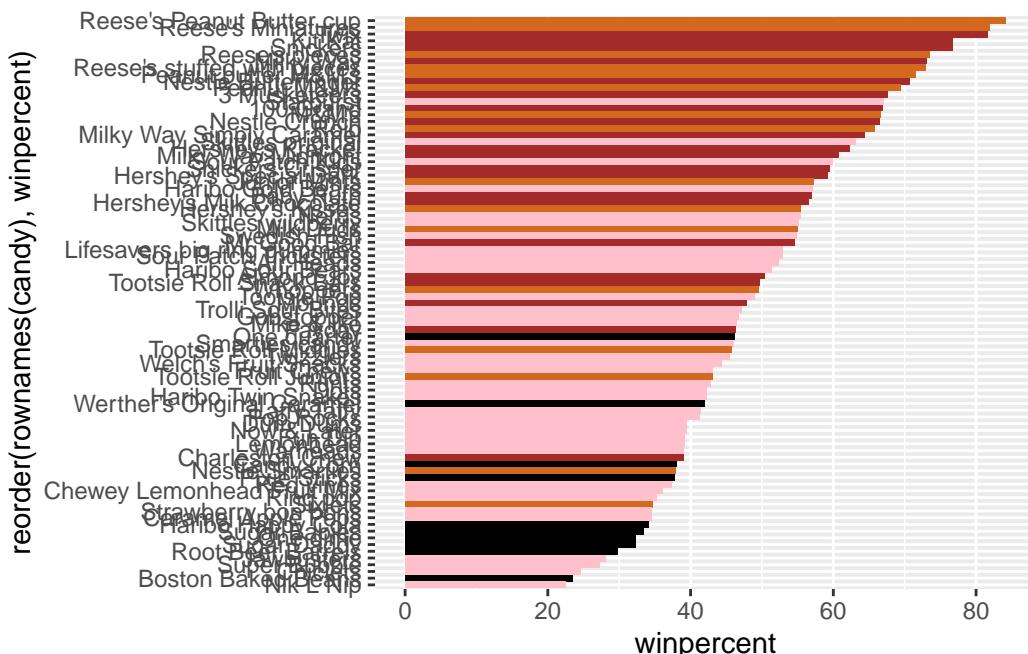
```
ggplot(candy, aes(winpercent, reorder(row.names(candy), winpercent))) +
  geom_col()
```



Time to add some useful color

```
my_cols=rep("black", nrow(candy))
my_cols[as.logical(candy$chocolate)] = "chocolate"
my_cols[as.logical(candy$bar)] = "brown"
my_cols[as.logical(candy$fruity)] = "pink"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=my_cols)
```



Q17. What is the worst ranked chocolate candy?

Sixlets are the worst-ranked chocolate candy.

Q18. What is the best ranked fruity candy?

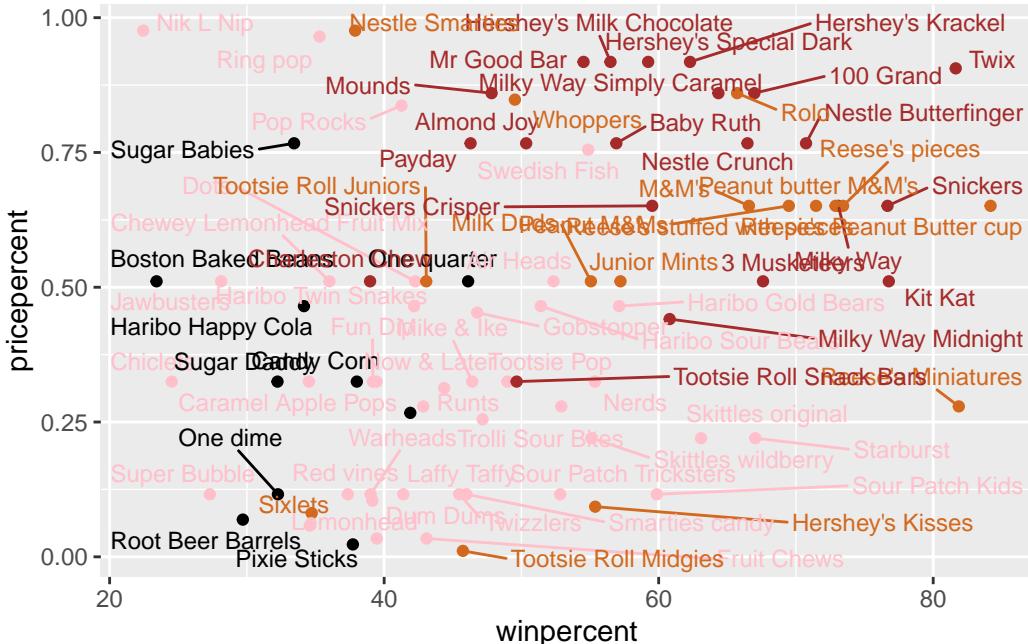
Starburst is the best-ranked fruity candy.

Taking a look at pricepercent

```
library(ggrepel)

# How about a plot of win vs price
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 17)
```

Warning: ggrepel: 4 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

Reese's Miniatures.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
price <- candy %>%
  arrange(desc(pricepercent))

head(price, 5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0	0	0	0
Nestle Smarties	1	0	0	0	0	0
Ring pop	0	1	0	0	0	0
Hershey's Krackel	1	0	0	0	0	0
Hershey's Milk Chocolate	1	0	0	0	0	0
	crisp	pedricewafer	hard	bar	pluribus	sugarpercent
Nik L Nip	0	0	0	1	1	0.197
Nestle Smarties	0	0	0	1	1	0.267
Ring pop	0	1	0	0	0	0.732

Hershey's Krackel	1	0	1	0	0.430
Hershey's Milk Chocolate	0	0	1	0	0.430
pricepercent winpercent					
Nik L Nip	0.976	22.44534			
Nestle Smarties	0.976	37.88719			
Ring pop	0.965	35.29076			
Hershey's Krackel	0.918	62.28448			
Hershey's Milk Chocolate	0.918	56.49050			

Nik L Nips is the most expensive AND least popular candy in this dataset.

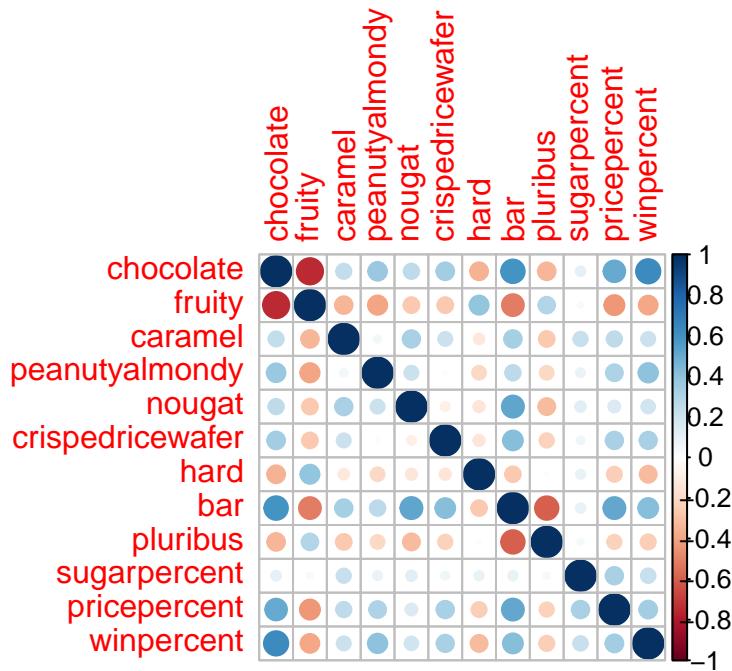
Exploring the correlation structure

Using the `corrplot` package (install first using `install.packages("corrplot")`):

```
library(corrplot)
```

`corrplot` 0.95 loaded

```
cij <- cor(candy)
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

Chocolate and fruity are anti-correlated. Additionally, pluribus and bar are also anti-correlated.

Q23. Similarly, what two variables are most positively correlated?

Chocolate and bar are the most positively correlated variables.

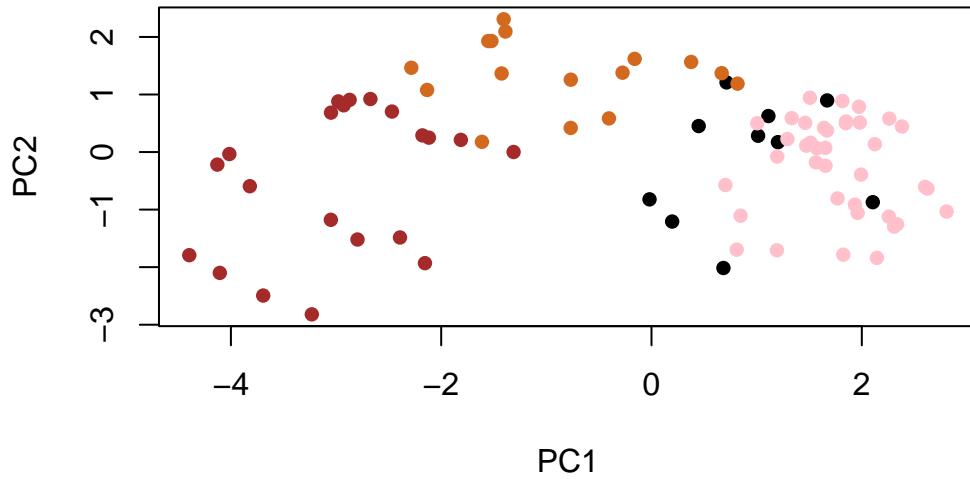
Principal Component Analysis

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369
	PC8	PC9	PC10	PC11	PC12		
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317		
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000		

```
plot(pca$x[, 1:2], col = my_cols, pch = 16)
```



Using `ggplot2` and `ggrepel` packages:

```
my_data <- cbind(candy, pca$x[,1:3])

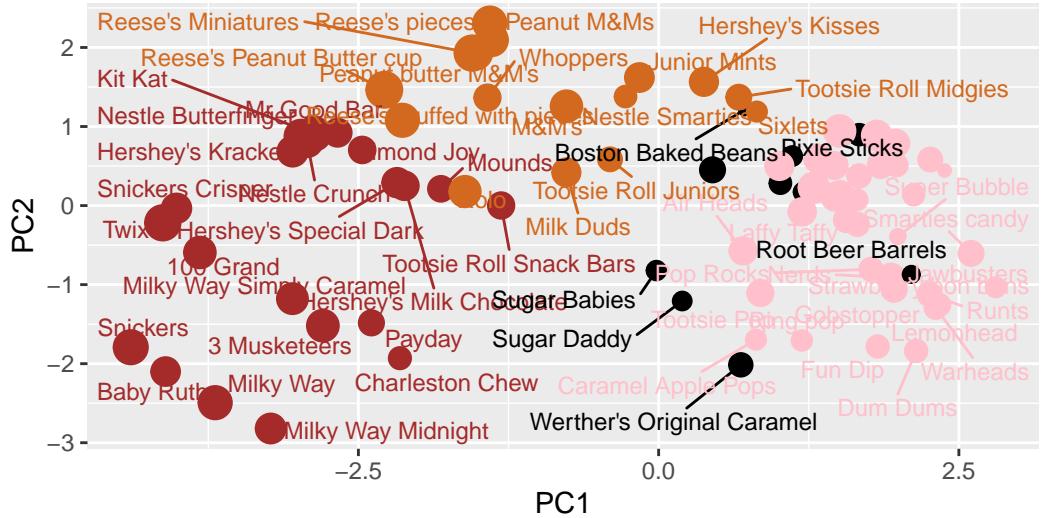
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=my_cols)

p + geom_text_repel(size=3.3, col=my_cols, max.overlaps = 17) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
       caption="Data from 538")
```

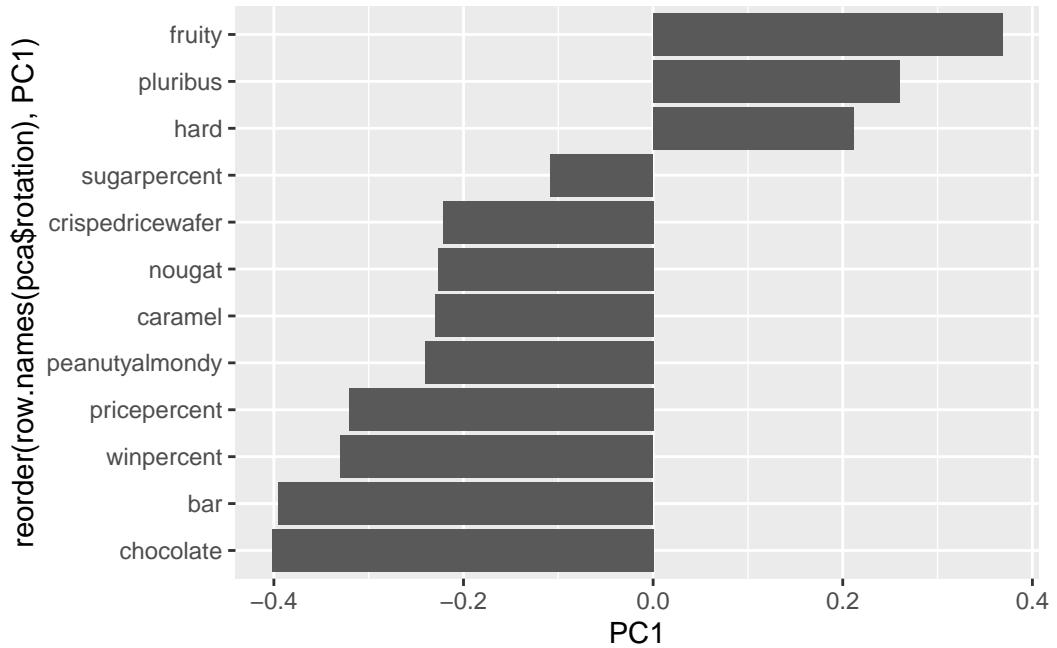
Warning: ggrepel: 25 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),



```
ggplot(pca$rotation) +
  aes(PC1, reorder(row.names(pca$rotation), PC1)) +
  geom_col()
```



Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

Fruity and pluribus are picked up strongly by PC1 in the positive direction. Given that I previously found the strongest anti-correlations between fruity & chocolate and pluribus & bar, this makes sense to me.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

The combination of characteristics that seem to make a “winning” candy is chocolate + bar.

From the barplot visualizations, you can qualitatively see that most of the chocolate/bar candies are at the top (where winpercent is highest).

From the correlation plot, you can see that winpercent is most positively correlated with chocolate, and in turn chocolate is also strongly positively correlated with bar.

From the PCA, you can see that candies that are chocolate bars make up their own cluster, and are picked up strongly by PC1 in the same direction as winpercent.

Optional extension questions

```
losers = candy[which(candy$winpercent < 50),]  
winners = candy[which(candy$winpercent >= 50),]
```

Q26. Are popular candies more expensive? In other words: is price significantly different between “winners” and “losers”? List both average values and a P-value along with your answer.

```
summary(losers$pricepercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0110	0.1160	0.3250	0.3744	0.5110	0.9760

```
summary(winners$pricepercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0930	0.4530	0.6510	0.5804	0.7670	0.9180

```
t.test(losers$pricepercent, winners$pricepercent)
```

Welch Two Sample t-test

```
data: losers$pricepercent and winners$pricepercent
t = -3.5653, df = 82.798, p-value = 0.0006068
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-0.32090727 -0.09107157
sample estimates:
mean of x mean of y
0.3743696 0.5803590
```

“Losers” have an average price percent of 37.44%, while “winners” have an average price percent of 58.04%. Conducting a t-test comparing the price percents of both groups returns a p-value < 0.05, suggesting that popular candies are significantly more expensive than “loser” candies.

Q27. Are candies with more sugar more likely to be popular? What is your interpretation of the means and P-value in this case?

```
non_sugary <- candy[which(candy$sugarpercent < 0.50),]
sugary <- candy[which(candy$sugarpercent > 0.50),]

t.test(non_sugary$winpercent, sugary$winpercent)
```

Welch Two Sample t-test

```
data: non_sugary$winpercent and sugary$winpercent
t = -2.1192, df = 76.967, p-value = 0.0373
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
-13.0167630 -0.4050308
sample estimates:
mean of x mean of y
47.23765 53.94854
```

Yes, candies with more sugar are likely to be more popular. Sugary candies have an average winpercent of 53.95%, while candies with less sugar have an average winpercent of 47.24%. Conducting a t-test reports a p-value < 0.05, suggesting that sugary candies are significantly more popular than non-sugary candies.